

Bootcamp AWS Clusters

To encourage collaboration across data science, solutions engineering, and project management, we have created clusters for each team. See the table below for the IP and domain addresses of these clusters.

	IP address	AWS Domain Address
Team 1 (Orange Lions)	54.197.41.71	ec2-54-197-41-71.compute-1.amazonaws.com
Team 2 (GAMBS)	54.208.153.19	ec2-54-208-153-19.compute-1.amazonaws.com
Team 3 ((Data Big Big)	54.167.220.218	ec2-54-167-220-218.compute-1.amazonaws.com
Team 4 (Rockers)	54.242.84.133	ec2-54-242-84-133.compute-1.amazonaws.com

All of these clusters have the following services pre-installed on them:

Service	Port	Login credentials
SSH	22	None (see SSH key below)
Resource Manager	8088	None
HDFS NameNode	50070	None
Spark 2.0.1	4040, 18080	None
Hive 2.1.0	CLI	None
Pig 0.16.0	CLI	None
Tez 0.8.4	CLI	None
Ganglia 3.7.2	IPaddress/ganglia	None
Zeppelin 0.6.2	8890	None
RStudio Server 1.0	8787	username: hadoop, password hadoop

To access these ports, you will need to join the Wireless-Guest or Wireless-Guest-5 networks. They are blocked on the Wireless-Employee network. You can find instructions for how to join the Wireless-Guest networks on the tables in the training rooms.

To log into your clusters, you will need to use ssh with a private key, which instructors will distribute to you. The key must have 400 (r--) permissions for ssh to accept it. The bash/gitbash command line you will need looks like the following:

```
ssh -i ~/keyname.pem hadoop@ec2-54-242-84-133.compute-1.amazonaws.com
```